

```
# Import libraries

import pandas as pd

import numpy as np

import matplotlib.pyplot as plt

import seaborn as sns

from sklearn.model_selection import train_test_split, GridSearchCV

from sklearn.preprocessing import StandardScaler

from sklearn.ensemble import RandomForestClassifier

from sklearn.metrics import classification_report, confusion_matrix, roc_curve, auc

import joblib


# Load the dataset

data = pd.read_csv("churn_data.csv")


# Display the first few rows

data.head()


# Check for missing values

print("Missing values per column:\n", data.isnull().sum())


# Fill or drop missing values (example: filling with mean)

data.fillna(data.mean(), inplace=True)


# Exploratory Data Analysis (EDA)

sns.countplot(x='Churn', data=data)

plt.title("Churn Distribution")

plt.show()
```

```
# Visualize correlations

corr_matrix = data.corr()

plt.figure(figsize=(12, 8))

sns.heatmap(corr_matrix, annot=True, fmt=".2f", cmap="coolwarm")

plt.title("Correlation Matrix")

plt.show()


# Feature engineering (convert categorical variables to dummy variables)

data = pd.get_dummies(data, drop_first=True)


# Define features and target

X = data.drop("Churn", axis=1)

y = data["Churn"]


# Train-test split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)


# Standardize features

scaler = StandardScaler()

X_train = scaler.fit_transform(X_train)

X_test = scaler.transform(X_test)


# Model training

rf = RandomForestClassifier(random_state=42)

rf.fit(X_train, y_train)
```

```
# Model evaluation

y_pred = rf.predict(X_test)

print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))

print("Classification Report:\n", classification_report(y_test, y_pred))


# ROC curve

y_pred_proba = rf.predict_proba(X_test)[:, 1]

fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)

roc_auc = auc(fpr, tpr)


plt.figure()

plt.plot(fpr, tpr, color='darkorange', lw=2, label=f"ROC curve (area = {roc_auc:.2f})")

plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')

plt.xlabel("False Positive Rate")

plt.ylabel("True Positive Rate")

plt.title("Receiver Operating Characteristic")

plt.legend(loc="lower right")

plt.show()


# Hyperparameter tuning using GridSearchCV

param_grid = {

    'n_estimators': [100, 200, 300],

    'max_depth': [None, 10, 20, 30],

    'min_samples_split': [2, 5, 10],

    'min_samples_leaf': [1, 2, 4]

}
```

```
grid_search = GridSearchCV(estimator=rf, param_grid=param_grid, cv=3, n_jobs=-1,
verbose=2)

grid_search.fit(X_train, y_train)


# Best parameters and evaluation

print("Best Parameters:\n", grid_search.best_params_)

best_model = grid_search.best_estimator_

y_pred_best = best_model.predict(X_test)

print("Optimized Classification Report:\n", classification_report(y_test, y_pred_best))


# Save the model and scaler

joblib.dump(best_model, "churn_prediction_model.pkl")

joblib.dump(scaler, "scaler.pkl")


# Example: Loading the model and making predictions

loaded_model = joblib.load("churn_prediction_model.pkl")

loaded_scaler = joblib.load("scaler.pkl")

new_data = pd.DataFrame([[25, 50, 3]], columns=X.columns) # Example new customer data

new_data_scaled = loaded_scaler.transform(new_data)

prediction = loaded_model.predict(new_data_scaled)

print("Prediction for new customer:", prediction)
```